

Name :Usman Manzoor

Roll no: P19-0068

Notebook settings

Hardware accelerator

GPU ?

To get the most out of Colab, avoid using a GPU unless you need one. [Learn more](#)

☐ Background execution

Want your notebook to keep running even after you close your browser? [Upgrade to Colab Pro+](#)

☐ Omit code cell output when saving this notebook

Cancel Save

```
!nvidia-smi
```

```
Fri Apr 29 19:13:45 2022
```

+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+									
NVIDIA-SMI		460.32.03		Driver Version: 460.32.03			CUDA Version: 11.2		
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+									
GPU	Name		Persistence-M		Bus-Id	Disp.A	Volatile Uncorr. ECC		
Fan	Temp	Perf	Pwr:Usage/Cap		Memory-Usage		GPU-Util	Compute M.	MIG M.
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+									
0	Tesla	T4	Off		00000000:00:04.0	Off	0		
N/A	37C	P8	9W / 70W		0MiB / 15109MiB		0%	Default	N/A
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+									
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+									
Processes:									
GPU	GI	CI	PID	Type	Process name			GPU Memory	
	ID	ID						Usage	
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+									
No running processes found									
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+									

```
!apt-get --purge remove cuda nvidia* libnvidia-*
!dpkg -l | grep cuda- | awk '{print $2}' | xargs -n1 dpkg --purge
!apt-get remove cuda-*
!apt autoremove
!apt-get update
```

```
Note, selecting 'nvidia-kernel-source-450-server' for glob 'nvidia*'
Note, selecting 'nvidia-440-utils' for glob 'nvidia*'
Note, selecting 'nvidia-compute-utils-440-server' for glob 'nvidia*'
Note, selecting 'nvidia-387-updates' for glob 'nvidia*'
Note, selecting 'nvidia-kernel-source-450-server' for glob 'nvidia*'
Note, selecting 'nvidia-current-diagnostic' for glob 'nvidia*'
Note, selecting 'nvidia-legacy-340xx-vdpau-driver' for glob 'nvidia*'
Note, selecting 'nvidia-349-updates' for glob 'nvidia*'
Note, selecting 'nvidia-utils-460-server' for glob 'nvidia*'
Note, selecting 'nvidia-opencl-icd-304-updates' for glob 'nvidia*'
Note, selecting 'nvidia-310-updates' for glob 'nvidia*'
```

```
[2] wget https://developer.nvidia.com/compute/cuda/9.2/Prod/local_installers/cuda-ubuntu1604-9-2-local_9.2.88-1_amd64 -O cuda-repo-ubuntu1604-9-2-local_9.2.88-1
!dpkg -i cuda-repo-ubuntu1604-9-2-local_9.2.88-1_amd64.deb
!apt-key add /var/cuda-repo-9-2-local/7fa2af80.pub
!apt-get update
!apt-get install cuda-9.2

Unpacking libxmu-headers (2:1.1.2-2) ...
Selecting previously unselected package libxmu-dev:amd64.
Preparing to unpack .../53-libxmu-dev_2%3a1.1.2-2_amd64.deb ...
Unpacking libxmu-dev:amd64 (2:1.1.2-2) ...
Selecting previously unselected package libxi-dev:amd64.
Preparing to unpack .../54-libxi-dev_2%3a1.7.9-1_amd64.deb ...
Unpacking libxi-dev:amd64 (2:1.7.9-1) ...
Selecting previously unselected package cuda-samples-9-2.
Preparing to unpack .../55-cuda-samples-9-2_9.2.88-1_amd64.deb ...
Unpacking cuda-samples-9-2 (9.2.88-1) ...
Selecting previously unselected package cuda-documentation-9-2.
Preparing to unpack .../56-cuda-documentation-9-2_9.2.88-1_amd64.deb ...
Unpacking cuda-documentation-9-2 (9.2.88-1) ...
Selecting previously unselected package cuda-libraries-dev-9-2.
Preparing to unpack .../57-cuda-libraries-dev-9-2_9.2.88-1_amd64.deb ...
Unpacking cuda-libraries-dev-9-2 (9.2.88-1) ...
```

```
!nvcc --version
```

```
nvcc: NVIDIA (R) Cuda compiler driver
Copyright (c) 2005-2020 NVIDIA Corporation
Built on Mon_Oct_12_20:09:46_PDT_2020
Cuda compilation tools, release 11.1, V11.1.105
Build cuda_11.1.TC455_06.29190527_0
```

```
!pip install git+https://github.com/andreinechaev/nvcc4jupyter.git
```

```
Collecting git+https://github.com/andreinechaev/nvcc4jupyter.git
  Cloning https://github.com/andreinechaev/nvcc4jupyter.git to /tmp/pip-req-build-ulfsr9m1
  Running command git clone -q https://github.com/andreinechaev/nvcc4jupyter.git /tmp/pip-req-build-ulfsr9m1
Building wheels for collected packages: NVCCPlugin
  Building wheel for NVCCPlugin (setup.py) ... done
  Created wheel for NVCCPlugin: filename=NVCCPlugin-0.0.2-py3-none-any.whl size=4306 sha256=8c1a3dffacf3fac116e74
  Stored in directory: /tmp/pip-ephem-wheel-cache-7h0fh8op/wheels/ca/33/8d/3c86eb85e97d2b6169d95c6e8f2c297fdec60
Successfully built NVCCPlugin
Installing collected packages: NVCCPlugin
Successfully installed NVCCPlugin-0.0.2
```

```
%load_ext nvcc_plugin
```

```
created output directory at /content/src  
Out bin /content/result.out
```

```
%%cu  
#include <stdio.h>  
#include <stdlib.h>  
__global__ void add(int *a, int *b, int *c) {  
    *c = *a + *b;  
}  
int main() {  
    int a, b, c;  
    // host copies of variables a, b & c  
    int *d_a, *d_b, *d_c;  
    // device copies of variables a, b & c  
    int size = sizeof(int);  
    // Allocate space for device copies of a, b, c  
    cudaMalloc((void **)&d_a, size);  
    cudaMalloc((void **)&d_b, size);  
    cudaMalloc((void **)&d_c, size);  
    // Setup input values  
    c = 0;  
    a = 3;  
    b = 5;  
    // Copy inputs to device  
    cudaMemcpy(d_a, &a, size, cudaMemcpyHostToDevice);  
    cudaMemcpy(d_b, &b, size, cudaMemcpyHostToDevice);  
    // Launch add() kernel on GPU  
    add<<<1,1>>>>(d_a, d_b, d_c);  
    // Copy result back to host  
    cudaError_t err = cudaMemcpy(&c, d_c, size, cudaMemcpyDeviceToHost);  
    if(err!=cudaSuccess) {  
        printf("CUDA error copying to Host: %s\n", cudaGetErrorString(err));  
    }  
    printf("result is %d\n",c);  
    // Cleanup  
    cudaFree(d_a);  
    cudaFree(d_b);  
    cudaFree(d_c);  
    return 0;  
}
```

```
result is 8
```